

MAY 27, 2020

Introduction to (Deep) Reinforcement Learning

Deep Learning Sessions Lisbon

Motivation

Humans learn autonomously how to make decisions

Why are humans so good at RL?

- People have prior experience
 - Humans learn from very few samples (*trial-and-error basis*)
- People have an existing representation of the world
 - Humans (subconsciously) use discrimination and generalization to identify and classify the world
- People are goal-oriented
 - Active behavior with sequential interactions
- Open questions about human behavior
 - Knowledge from evolution, culture, experience, and so on ...

Mapping questions

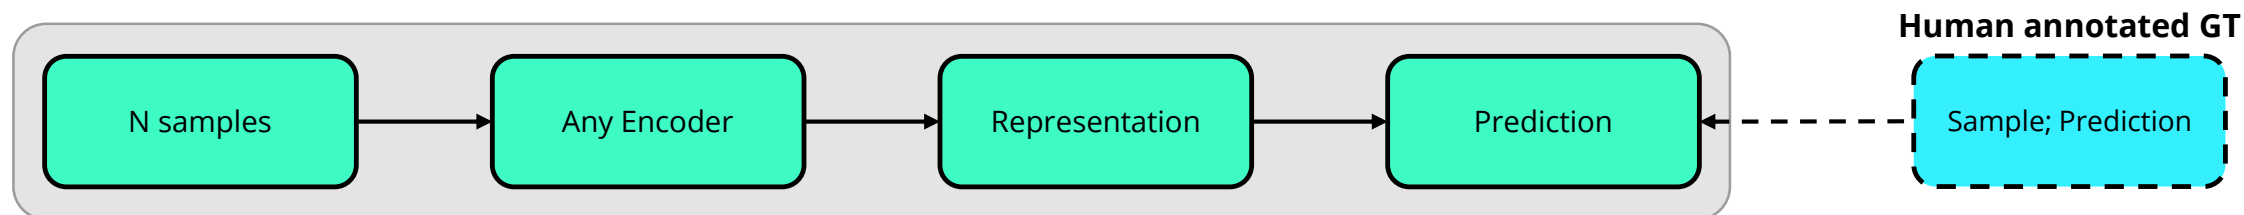
- Is world model learning an essential step to be learned before or in parallel with policy model?
- Can we learn a representation under which RL solve new tasks under very small amount of experience?

Type of Learning

- Optimization
- Generalization


Supervised Learning

“Teach by example” – design of the annotations




Input

x



Network

$$y = f(x; \theta)$$



Output

y



Data

$\{(x, y)_i\}$




Type of Learning

- Optimization
- Generalization
- Delayed consequences

Imitation Learning

behavioral cloning (supervised learning) for sake of simplicity

“Teach by expert” – design to mimic




Input

s_t



Network

$$a_t = \pi(s_t; \theta)$$



Output

a_t



Data

$\{(s_t, a_t^*)\}$




Type of Learning

- Optimization
- Generalization
- Delayed consequences
- Exploration

Reinforcement Learning

“Teach by experience” – design of the world




Input

s_t



Network

$$a_t = \pi(s_t; \theta)$$



Output

a_t





Data

$$\{(s_t, a_t, r_t, s_{t+1})\}$$




Applications



Interaction Loop

Agent interacts with an environment by following a policy



- The agent contains:

- Agent state
- Policy
- Value function (*probably*)
- Model (*optionally*)

- At each step t the agent:


- Receives observation o_t (and reward r_t)
- Executes a_t

- The environment:

- Receives action a_t
- Emits observation o_{t+1} (and reward r_{t+1})

- Aspire learning to make decisions from interaction ...

- ... time;
- ... long-term consequence of actions;
- ... actively gathering experience;
- ... predict the future;
- ... deal with uncertainty




- **Optimization**
- **Generalization**
- **Delayed consequences**
- **Exploration**

Final goal: to learn an optimal policy which maximizes the long-term cumulative rewards

Example #1

Robot labyrinth



Actions

{UP, DOWN, LEFT, RIGHT}

States

Cells

Questions & Options

What's the strategy to achieve max reward?

- Learn the model and plan
- Learn the value of (action, state) pairs and act greedy/non-greedy
- Learn the policy directly while sampling from it

Deterministic model of the world

Policy:

- **Shortest path**

Stochastic model of the world

Policy:

Reward for each step: [-2, -0.04]

- **Shortest path**, however the structure changes; the higher negative reward for each step, the urgency of the agent increases, and vice-versa

Reward for each step: +

- **Longest path**

Notation

States & Observations

- A *state* \mathbf{s} is a complete description of the state – no hidden information
- An *observation* \mathbf{o} is a partial description of a state – omit information
- Both agent and environment may have an internal state

Environment State

- It is the environment's internal state
- *Fully observed* – agent is able to observe the complete state of the environment
- *Partially observed* – agent only see a partial observation

Agent State

- Actions depend on the state
- A history is a sequence of

$$\mathcal{H}_t = o_0, a_0, r_1, o_1, \dots, o_{t-1}, a_{t-1}, r_t, o_t$$

- The history can be used to construct an agent state s_t

Notation

Trajectories

- A trajectory τ is a sequence of states and actions in the world

$$\tau = (s_0, a_0, s_1, a_1, \dots)$$

- Initial state s_0 randomly sampled from the start-state distribution

$$s_0 \sim \rho_0(\cdot)$$

- State transitions are governed by the laws of the environment, and depend on only the most recent action, \mathbf{a}_t :

$$\textit{deterministic} \rightarrow s_{t+1} = f(s_t, a_t)$$

$$\textit{stochastic} \rightarrow s_{t+1} \sim P(\cdot | s_t, a_t)$$

Notation

Reward & Return

- A reward r_t is a scalar feedback signal
- Indicates how well the agent is doing at step t
- Depends on the current state of the environment, the action just taken, and the next state of the world

$$r_t = R(s_t, a_t, s_{t+1})$$

- The goal of the agent is to maximize some notion of cumulative reward over a trajectory, so-called **return** $R(\tau)$
 - *Finite-horizon undiscounted return*: sum of rewards obtained in a fixed window of steps

$$R(\tau) = \sum_{t=0}^T r_t$$

- *Infinite-horizon discounted return*: sum of all rewards ever obtained by the agent, but discounted by how far off in the future they are obtained

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t, \text{ where } \gamma \in [0, 1]$$

Notation

Value Function

- The value of state (*or state-action pair*) is the expected cumulative reward, from a state s , and then act according to a particular policy π forever after

$$V(s) = \mathbb{E}[R(\tau) \mid s_t = s]$$

- The goal is to maximize value by picking suitable actions
- Rewards and values define desirability of a state or action
- There are four main functions:

- On-policy value function

$$V^\pi(s) = \mathbb{E}_{\tau \sim \pi}[R(\tau) \mid s_t = s]$$

- On-policy action-value function

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi}[R(\tau) \mid s_t = s, a_t = a]$$

- Optimal value function

$$V^*(s) = \max_{\pi} \mathbb{E}_{\tau \sim \pi}[R(\tau) \mid s_t = s]$$

- Optimal action-value function

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{\tau \sim \pi}[R(\tau) \mid s_t = s, a_t = a]$$

Notation


Policy

- Actions may have long term consequences
- Reward may be delayed
- A mapping from states to actions is called a policy – rule used by an agent to decide what actions to take
- *Parameterized policies*, whose outputs are computable functions that depend on a set of parameters
 - *Deterministic policy* - $\mu: a_t = \mu_\theta(s_t) \equiv \pi(s) = a$
 - *Stochastic policy* - $\pi: a_t \sim \pi_\theta(\cdot|s_t) \equiv \pi(a|s) = p(a_t = a|s_t = s)$

Notation

Model

- Agent's representation of the environment



- A model predicts what the environment will do next


$$\text{transition function} \rightarrow P(s', r | s, a) = p(s_{t+1} = s', r_{t+1} = r | s_t = s, a_t = a)$$

$$\text{next state} \rightarrow P(s' | s, a) = p(s_{t+1} = s' | s_t = s, a_t = a)$$

$$\text{next reward} \rightarrow R(s, a, s') \approx \mathbb{E}[r_{t+1} | s_t = s, a_t = a, s_{t+1} = s']$$

- A model does not immediately provide a good policy – planning is still needed
- Stochastic (generative) models can be used

Formalizing RL Interaction



- Which action is better or worse?
- $r(s, a)$: reward function, tell us which states and actions are better
- The mathematical formulation of the agent-environment interaction is called a Markov decision process (MDP)

$$s, a, R(s, a, s'), P(s' | s, a)$$

- Suppose the agent sees the full environment state (*fully observable*) $s_t = o_t = \text{environment state}$
 - Then the agent is in a MDP
- Suppose the agent gets partial information (*partially observable*)
 - Then it is a partially observable MDP (POMDP), where the observation is not Markov, the environment state can still be Markov, but the agent does not know it

MDP Context

Definition

- *“The future is independent of the past given the present”*
- Considering a sequence of random states, $\{S_t\}_{t \in \mathbb{N}}$, indexed by time, a state s has the *Markov* property when for states $\forall s' \in S$ and all rewards $r \in \mathbb{R}$

$$p(r_{t+1} = r, s_{t+1} = s' | s_t = s) = p(r_{t+1} = r, s_{t+1} = s' | s_1, \dots, s_{t-1}, s_t = s), \text{ for all possible } s_1, \dots, s_{t-1}$$



Markov property

- A MDP is a tuple (S, A, P, R, γ) , where
 - S is the set of all possible states
 - A is the set of all possible actions
 - $P(s' | s, a)$ is the conditional probability of next state s' , given a state s and action a – defines the dynamics
 - $R(s, a, s')$ is the joint probability of how much reward – defines the reward
 - $\gamma \in [0,1]$ is a discount factor that gives more importance to early rewards – rewards and discount define the goals

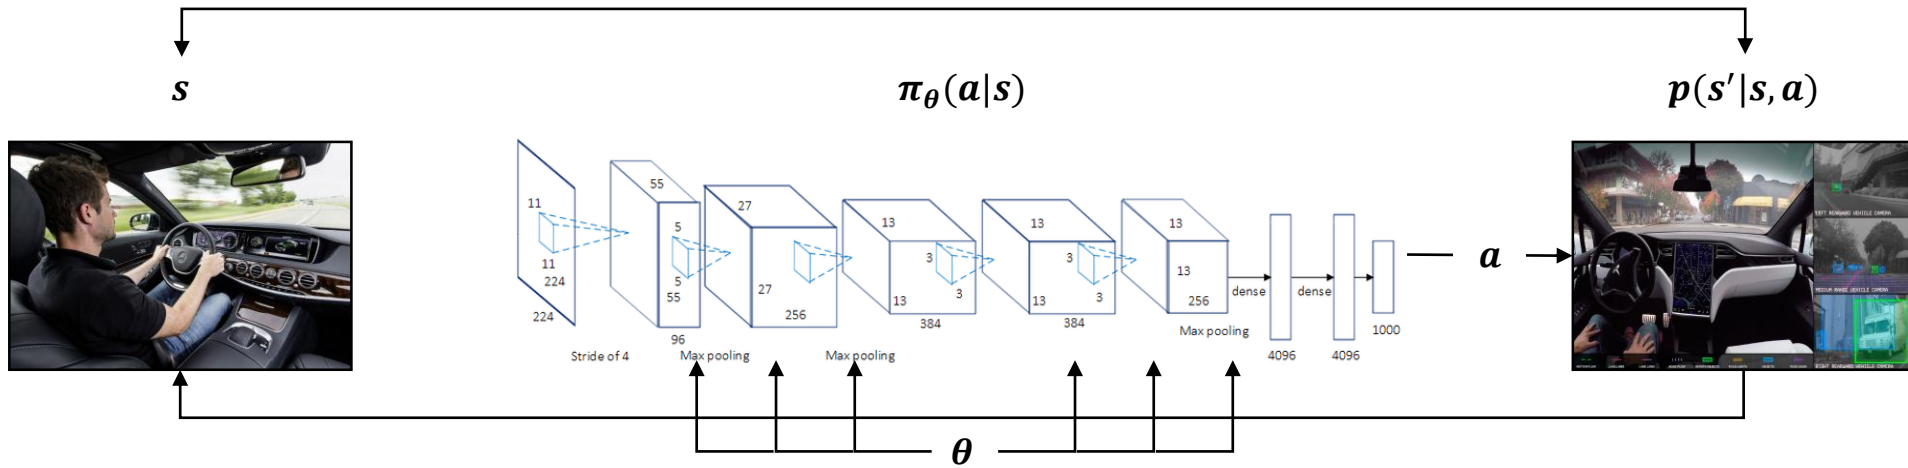
RL as MDP

Definition

- $\text{MDP} \rightarrow M = \{S, A, P, R, \gamma\}$
 - S – state space $s \in S$ (discrete or continuous)
 - A – action space $a \in A$ (discrete or continuous)
 - P – transition operator (tensor)
 - R – reward function $R: S \times A \rightarrow \mathbb{R}$
 - $\gamma \in [0,1]$ is a discount factor
- $\text{POMDP} \rightarrow M = \{S, A, O, T, \mathcal{E}, r, \gamma\}$
 - O – observation space $o \in O$ (discrete or continuous)
 - \mathcal{E} – emission probability $p(o_t|s_t)$




(Deep) RL Problem



- The goal is to select a policy which maximizes expected return when the agent acts according to it
- Assuming that both state transitions and policy are stochastic, the probability of a T -step trajectory is

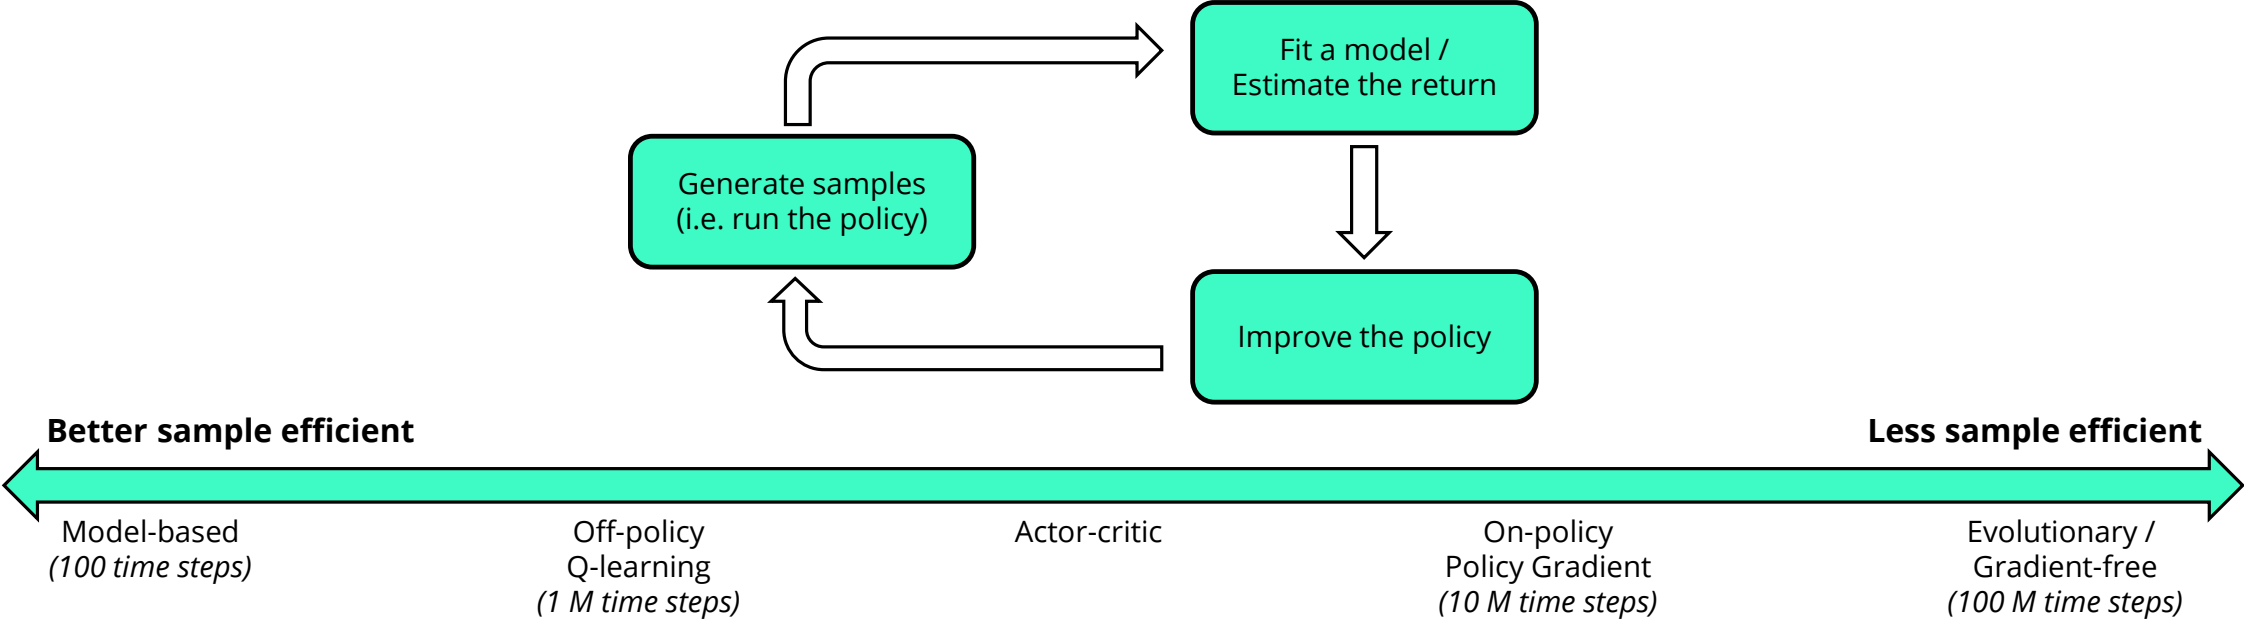
$$P(\tau|\pi) = \rho_0(s_0) \prod_{t=0}^{T-1} \underbrace{P(s_{t+1}|s_t, a_t) \pi(a_t|s_t)}_{\text{Markov chain on } (s, a)}$$

- The expected return is given by $J(\pi) = E_{\tau \sim \pi}[R(\tau)]$
- The optimal policy is given by $\pi^* = \underset{\pi}{\operatorname{argmax}} J(\pi)$



Types of RL

Anatomy of RL algorithms



Model-based

Estimate the transition model to
 Use it for planning (no explicit policy)
 Use it to improve a policy
 Update model and re-plan often

Value-based

Estimate value function or Q-function of the optimal policy (no explicit policy)
 Act by using best action in state
 Exploration is a necessary add-on

Actor-critic

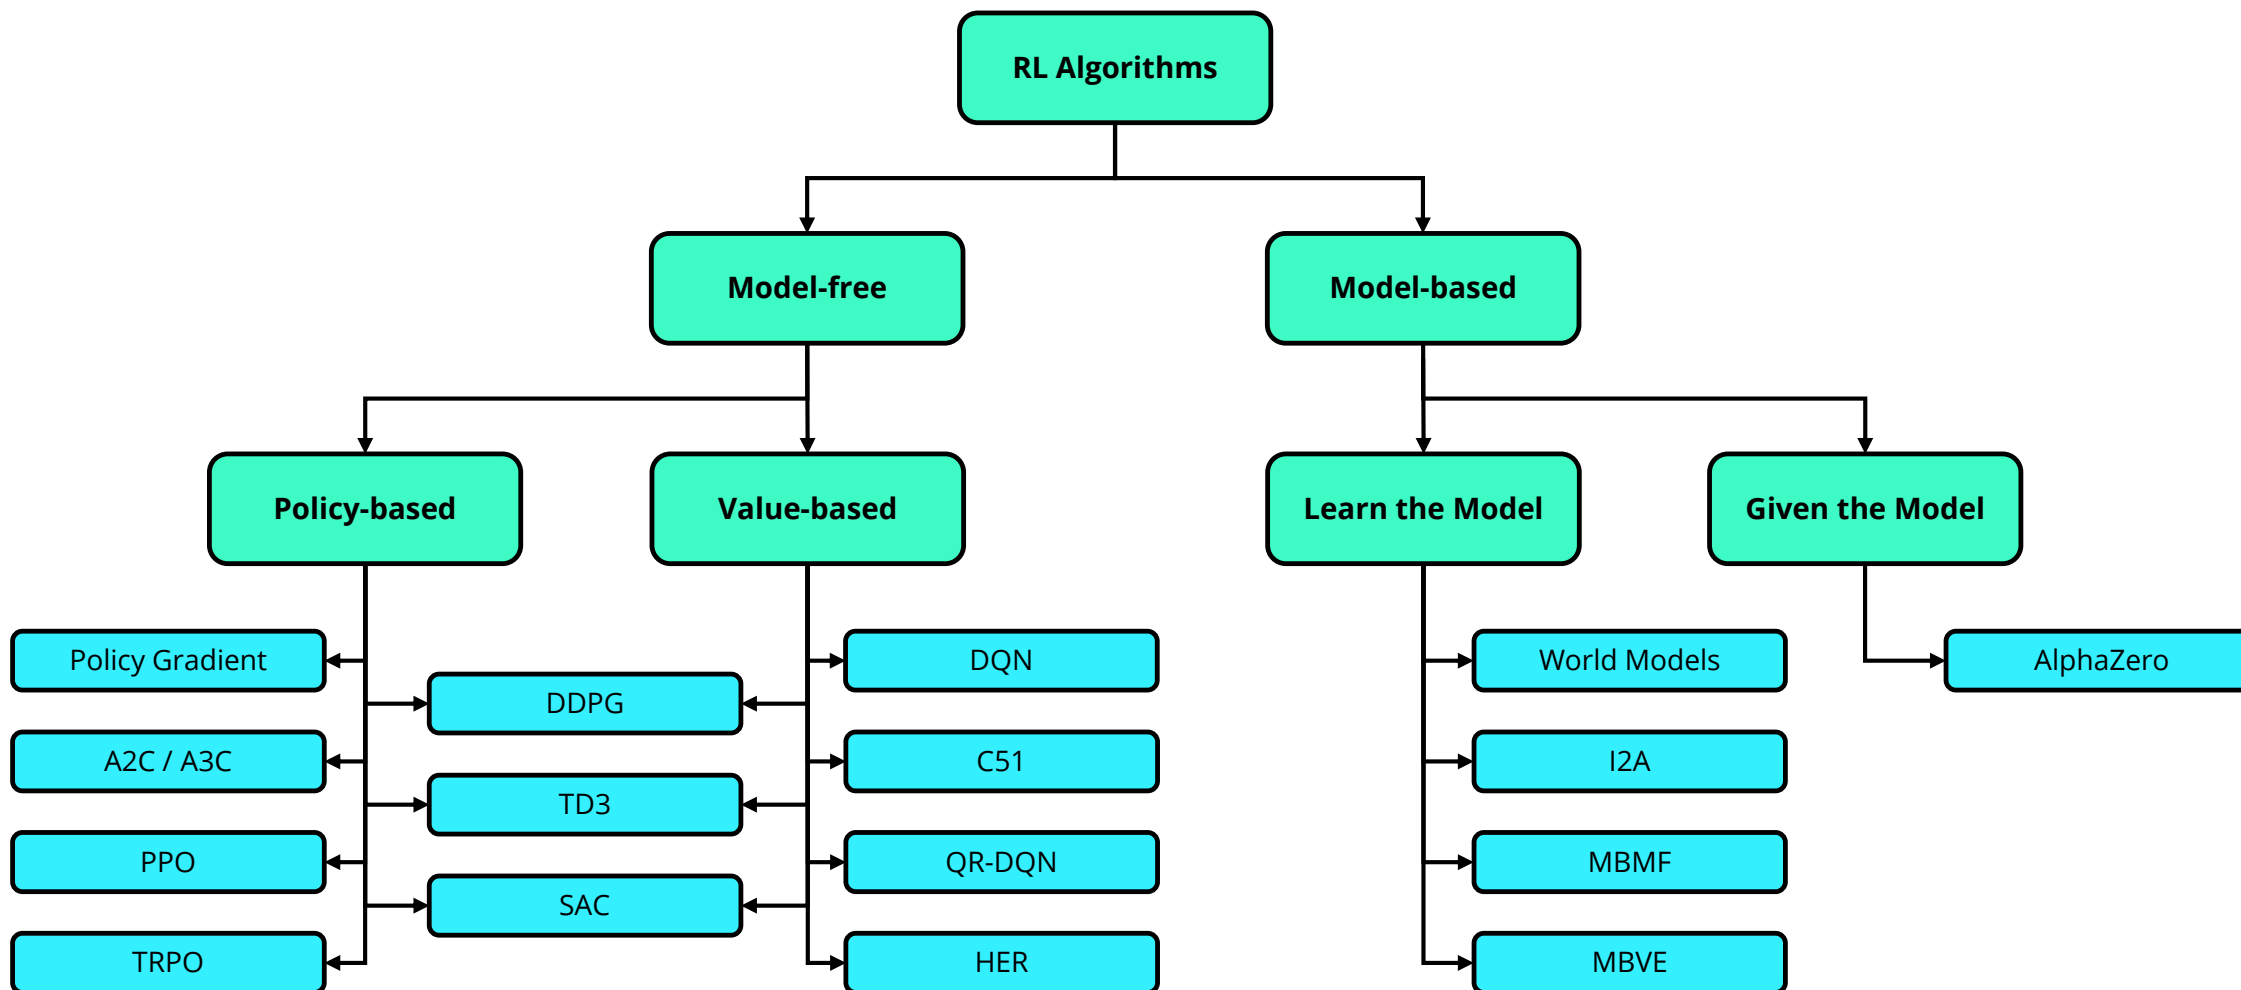
Estimate value function or Q-function of the current policy
 Use it to improve policy

Policy-based

Learn the stochastic policy function that maps state-to-action
 Act by sampling policy
 Exploration is baked-in

Taxonomy of RL

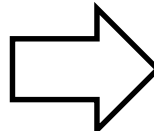
Different methods of RL



From: Spinningup - Openai

Challenges of (Deep) RL

Learning the components of an agent


- All components are functions
 - Policies map states to actions
 - Value functions map states to values
 - Models map states to states and/or rewards
 - State updates map states and observations to new states
 - These functions can be represented as neural networks
- challenges**
- 
- Learning
 - The environment is initially unknown
 - The agent interacts with the environment
 - Planning
 - A model of the environment is given
 - The agent plans in this model
 - Prediction
 - Evaluate the future for a given policy
 - Control
 - Optimize the future, finding the best policy

Example #2

Cartpole

Observation

Num	Observation	Min	Max
0	Cart Position	-2.4	2.4
1	Cart Velocity	$-\text{Inf}$	Inf
2	Pole Angle	-41.8°	41.8°
3	Pole Velocity At Tip	$-\text{Inf}$	Inf



Actions

Num	Action
0	Push cart to the left
1	Push cart to the right

Description

A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track. The pendulum starts upright, and the goal is to prevent it from falling over by increasing and reducing the cart's velocity.

Reward

Reward is 1 for every step taken, including the termination step. The threshold is 475 for v1.

Starting State

All observations are assigned a uniform random value between ± 0.05 .

Episode Termination

- Pole Angle is more than $\pm 12^\circ$
- Cart Position is more than ± 2.4 (center of the cart reaches the edge of the display)
- Episode length is greater than 200 (500 for v1).

Bellman Equations

Goal: solve MDP by finding the optimal policy and value functions

- Bellman equations decompose the value function

$$v = \underset{\substack{\uparrow \\ \text{immediate} \\ \text{reward}}}{r} + \underbrace{\gamma P^\pi v}_{\substack{\text{Discounted sum} \\ \text{of future} \\ \text{rewards}}}$$

- The Bellman expectation equations are given by

$$V^\pi(s) = \sum_{a \in A} \pi(a|s) (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^\pi(s'))$$

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') Q^\pi(s', a')$$

- The optimal values are given by

$$V^*(s) = \max_{a \in A} (R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^*(s'))$$

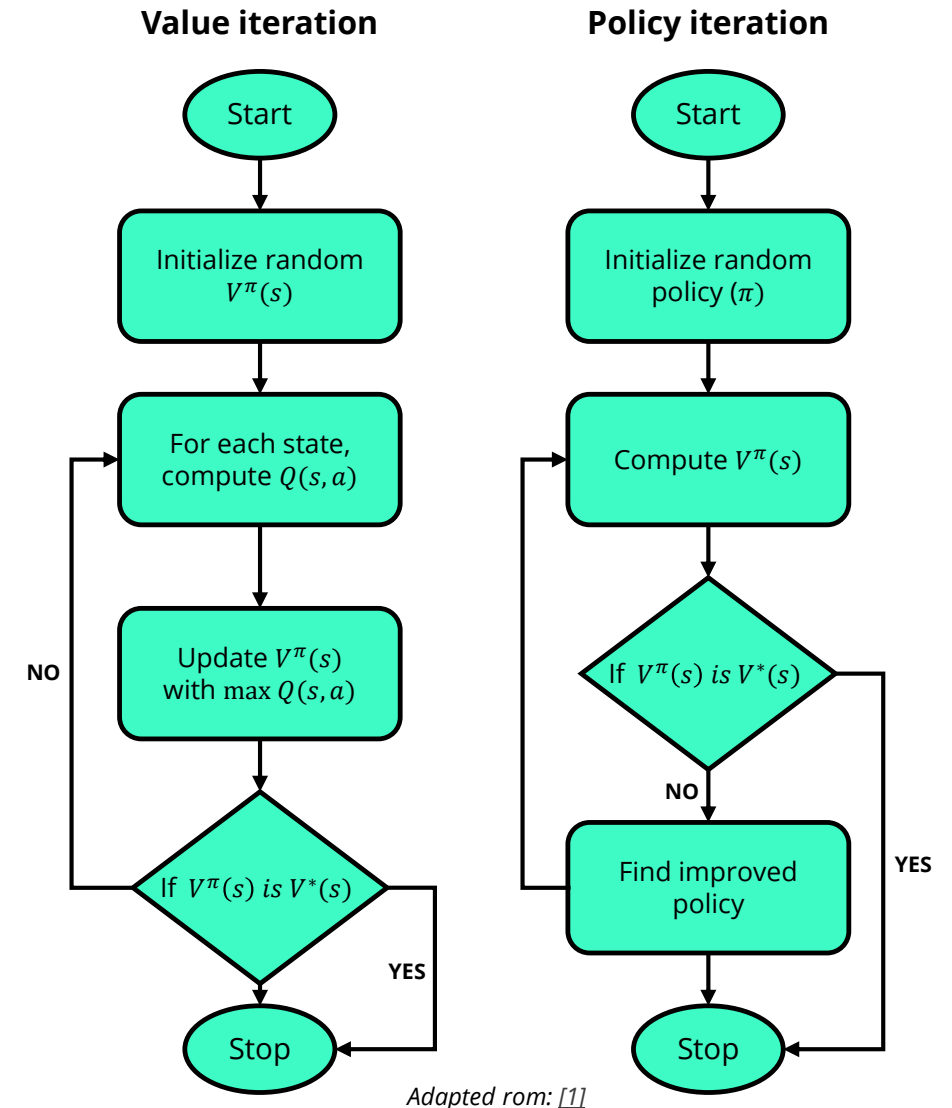
$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \max_{a' \in A} Q^*(s', a')$$

- Optimal policy

$$\pi^*(s) = \operatorname{argmax}_\pi V^\pi(s)$$

- The relative advantage of an action (compared to others on average) is

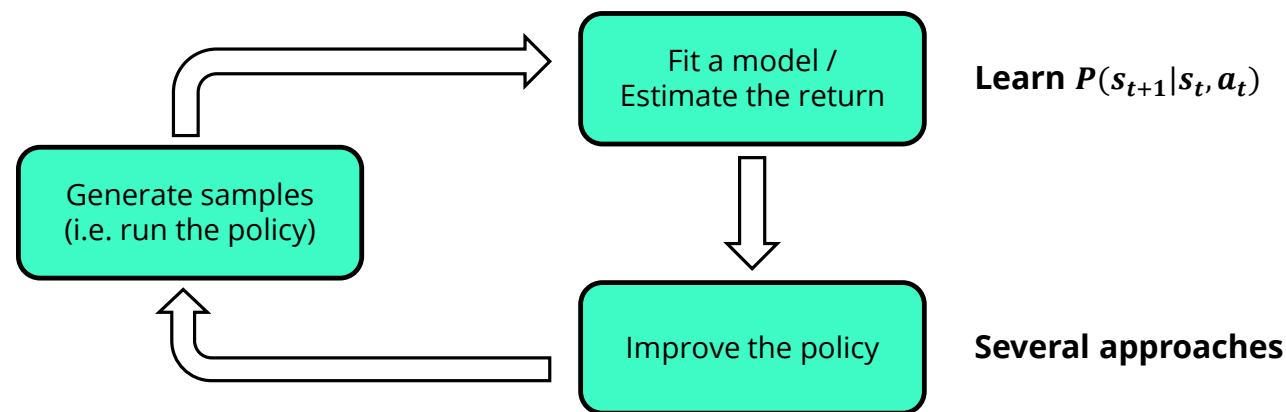
$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$



Model-based RL

You own the (real or simulated) representation of the environment you are in

Have a representation of
 $M = \{S, A, P, R, \gamma\}$



Goal: follow Bellman equations to iteratively evaluate value functions and improve policy

- Planning problem: policy iteration is generally more efficient than enumeration

Set $i = 0$


Initialize $\pi_0(s)$ randomly for all states s

While $i == 0$ or $\|\pi_i - \pi_{i-1}\|_1 > 0$:

$V^{\pi_i} \leftarrow$ MDP V function **policy evaluation** of π_i

$\pi_{i+1} \leftarrow$ **Policy improvement**

$i = i + 1$



- Policy evaluation

$$V^{\pi_i}(s) = R(s, \pi_i(s)) + \gamma \sum_{s' \in S} P(s'|s, \pi_i(s)) V^{\pi_{i-1}}(s')$$

- Policy improvement

$$Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^{\pi_i}(s')$$

$$\pi_{i+1} = \operatorname{argmax}_a Q^{\pi_i}(s, a)$$

- Monotonic improvement in policy

$$V^{\pi_1} \geq V^{\pi_2}: V^{\pi_1}(s) \geq V^{\pi_2}(s), \forall s \in S$$

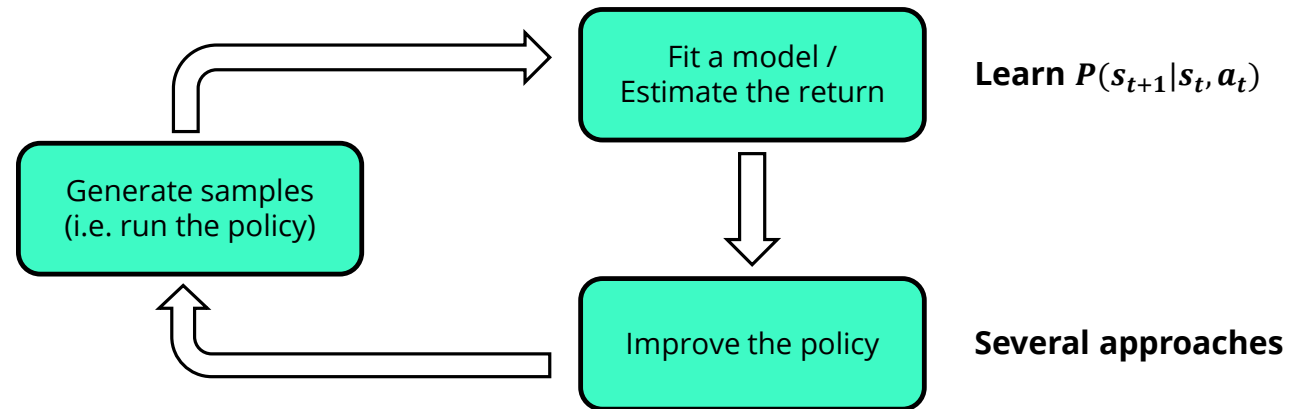
$$Q_{\pi}(s, \pi'(s)) = Q_{\pi}(s, \operatorname{argmax}_a Q_{\pi}(s, a))$$

$$= \max_a Q_{\pi}(s, a) \geq Q_{\pi}(s, \pi(s)) = V_{\pi}(s)$$

Model-based RL

You own the (real or simulated) representation of the environment you are in

Have a representation of
 $M = \{S, A, P, R, \gamma\}$



Analytic gradient computation

- Assumptions about the form of the dynamics and cost function
 - LQR framework, Receding-horizon control, other analytic gradient-based approaches used for policy improvement

Sampling-based planning

- Generate sampling distributions of action sequences (*continuous space*) or to search over tree structures (*discrete space*)
 - Random shooting, cross-entropy method, path integral optimal control, etc

Model-based data generation

- Increase size of training set for policy optimization
 - Dyna, iLOG, meta-learning, etc

Model-free RL

You don't know the dynamics and reward models

No knowledge of MDP is required, just samples

$$V^\pi(s) = \sum_{a \in A} \pi(a|s)(R(s, a) + \gamma \sum_{s' \in S} \boxed{P(s'|s, a)} V^\pi(s'))$$

Unknown!!

Value-based


- Learn value function
- Implicit policy (e.g. $\epsilon - greedy$)

Policy-based

- No value function
- Learnt policy

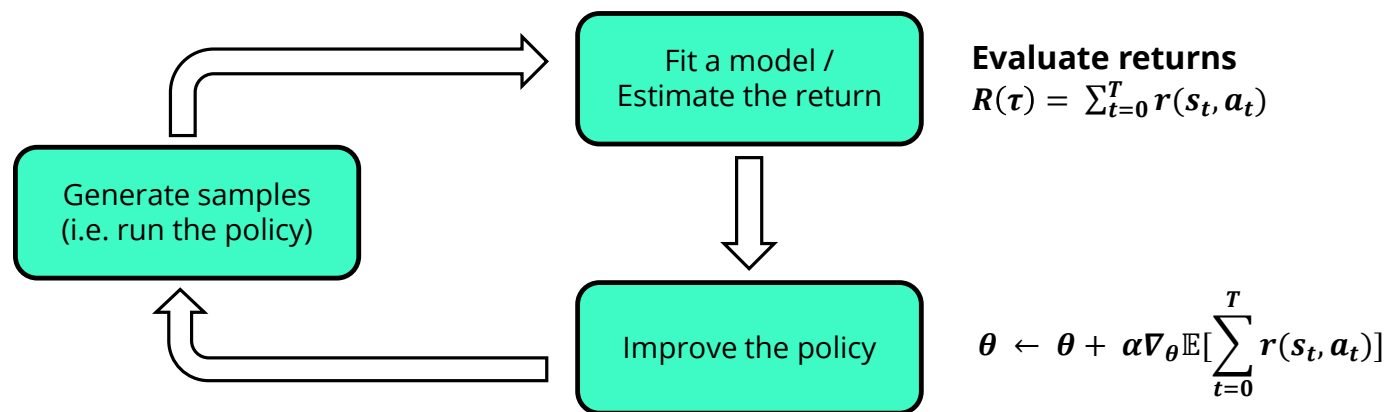
Actor-Critic

- Learnt value function
- Learnt policy



Model-free RL

Direct policy gradients



Goal: given policy $\pi_{\theta}(s, a)$ with parameters θ , find best θ

- The quality of a policy π_{θ} can be measured in
 - Episodic environments by the start value

$$J_0(\theta) = V^{\pi_{\theta}}(s_0)$$
 - Continuing environments by the average value

$$J_{avV}(\theta) = \sum_{s \in S} \mu^{\pi_{\theta}}(s) V^{\pi_{\theta}}(s)$$

or by the average reward per time-step

$$J_{avR}(\theta) = \sum_{s \in S} \mu^{\pi_{\theta}}(s) \sum_{a \in A} \pi(a|s; \theta) \sum_{r \in R} p(r|s, a) r$$

Optimization: find θ that maximizes $J(\theta)$

- No gradient-based
 - Hill climbing
 - Genetic algorithms
- Gradient based (stochastic gradient ascent)
 - Search for a local maximum in $J(\theta)$ by ascending the gradient of the policy, w.r.t parameters θ

$$\Delta \theta = \alpha \nabla_{\theta} J(\theta)$$

where $\nabla_{\theta} J(\theta)$ is the policy gradient approximated by


$$\nabla_{\theta} J(\theta) \propto \sum_{s \in S} \mu(s) \sum_{a \in A} Q_{\pi}(s, a) \pi(a|s; \theta) = \mathbb{E}^{\pi_{\theta}} [\nabla \log \pi(a|s, \theta) Q_{\pi}(s, a)]$$

Policy Evaluation Methods

Learning methods


Monte-Carlo (MC)

$$V^\pi(s_t) = V^\pi(s_t) + \alpha(R(\tau)_t - V^\pi(s_t))$$




Temporal-Difference (TD)

$$V^\pi(s_t) = V^\pi(s_t) + \alpha(R_{t+1} + \gamma V^\pi(s_{t+1}) - V^\pi(s_t))$$



Dynamic Programming (DP)

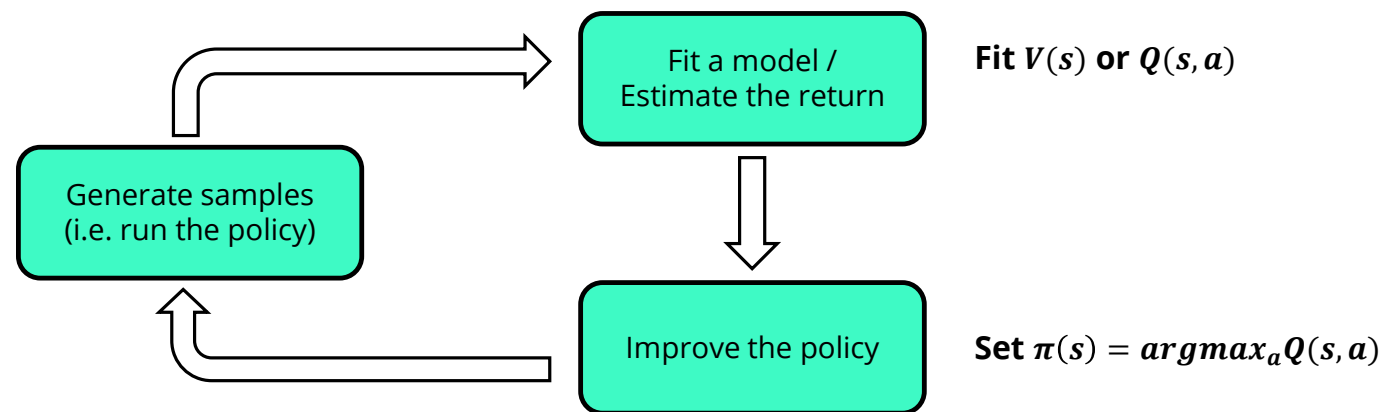
$$V^\pi(s_t) = [R_{t+1} + \gamma V^\pi(s_{t+1})]$$



From: David Silver's RL course

Model-free RL

Value functions



Goal: measure the (approximated) goodness of a state or how regarding a state or an action is by predicting the future reward

- The objective function is based on the Bellman equation

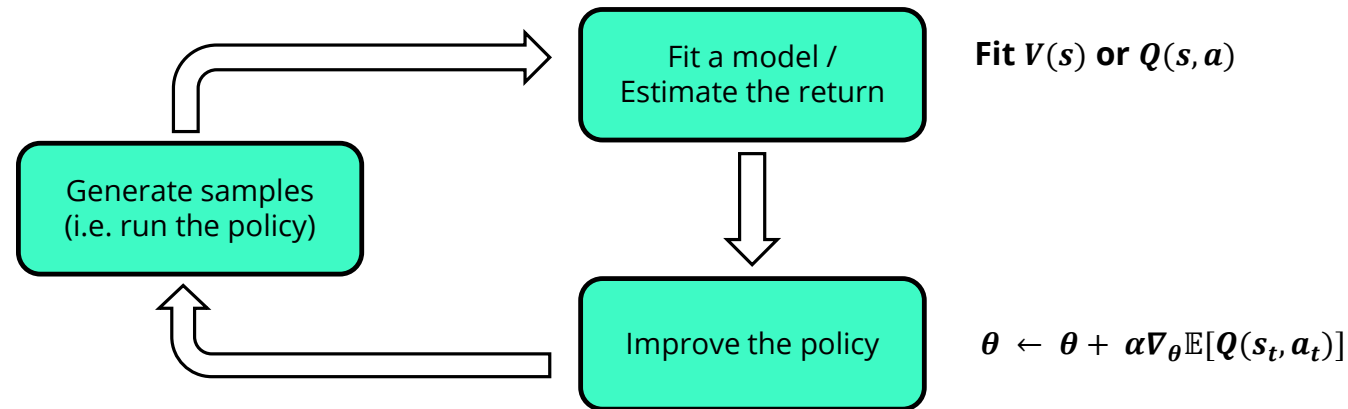
$$\mathbf{v} = \mathbf{r} + \gamma \mathbf{P}^\pi \mathbf{v}$$

- Monte-Carlo-based (MC) approach
 - Computes the observed mean return as an approximation of the expected return
 - Learns from complete episodes
- Temporal-Difference-based (TD) approach
 - Uses bootstrapping to update $V(s)$ towards an estimated return target
 - Learns from incomplete episodes

- Problems of value functions for large MDPs:
 - There are too many states and/or actions to store
 - It is too slow to learn the value of each individual state
 - Individual states are often not fully observable
- Function approximation:
 - Generalize from seen states to unseen states
 - Update parameter θ using TD or MC learning
 - For a partial observable environment, the agent state or a learning state update function can be used
 - ANN, decision tree, NN, fourier / wavelet bases, etc

Model-free RL

Actor-critic: value functions + policy gradients



Goal: learn value function in addition to the policy

Critic : update value function parameters w by (n-step) TD

Actor: update policy parameters θ by policy gradient

- Considering the policy-based gradient

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \boxed{R(\tau^i)} \longrightarrow R(\tau^i) - b(\tau^i)$$

- Improving high variance of gradients and cumulative reward of 0

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{t=0}^T \nabla_{\theta} \underbrace{\log \pi_{\theta}(a_t | s_t)}_{\text{Actor}} \underbrace{(Q(s_t, a_t) - V_{\phi}(s_t))}_{\text{Critic}}$$

advantage value

- The **critic** estimates the value function, could be the action-value (Q value) or state-value (V value)
- The **actor** updates the policy distribution in the direction suggested by the critic (through policy gradients)

Algorithm 1 Q Actor Critic

Initialize parameters s, θ, w and learning rates $\alpha_{\theta}, \alpha_w$; sample $a \sim \pi_{\theta}(a|s)$.
for $t = 1 \dots T$: **do**
 Sample reward $r_t \sim R(s, a)$ and next state $s' \sim P(s'|s, a)$
 Then sample the next action $a' \sim \pi_{\theta}(a'|s')$
 Update the policy parameters: $\theta \leftarrow \theta + \alpha_{\theta} Q_w(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)$; Compute the correction (TD error) for action-value at time t :
 $\delta_t = r_t + \gamma Q_w(s', a') - Q_w(s, a)$
 and use it to update the parameters of Q function:
 $w \leftarrow w + \alpha_w \delta_t \nabla_w Q_w(s, a)$
 Move to $a \leftarrow a'$ and $s \leftarrow s'$
end for

Exploration vs. Exploitation

We learn by trial and error

Policy

- Online decision-making involves a fundamental choice to discover a good policy:
 - Exploration – increase knowledge
 - Exploitation – maximize performance based on current knowledge
- The best long-term strategy may involve short-term sacrifices – enough information to make the best overall decision
- ϵ – *greedy* policy:
 - Greedy can lock onto a suboptimal action forever
 - ϵ – *greedy* continues to explore forever
 - Don't know anything about the environment at the beginning
 - Need to try all actions to find the optimal one
 - With probability $(1 - \epsilon)$ select $a = \operatorname{argmax}_{a \in A} Q_t(a)$
 - With probability ϵ select a random action

Off-policy RL

Q-Learning, an off-policy TD control


Goal: how to act optimally in controlled Markovian domains, by successively improving the evaluations of the quality of particular actions at particular states

- The optimization is almost always off-policy and the corresponding policy is obtained via the connection between Q^* and π^*

$$a(s) = \operatorname{argmax}_a Q_\theta(s, a)$$

- Independent of the policy being followed
- $\epsilon - greedy$ is commonly applied
- Only requirement: keep updating each pair (s, a)
- Classes of function approximation
 - Tabular – a table with an entry for each MDP state
 - State aggregation – partition environment states
 - Linear function approximation – fixed features/kernel
 - Differentiable (non-linear) function approximation - NNs

↓
Deep Q-Network



	A1	A2	A3	A4
S1	+1	+2	-1	0
S2	+2	0	+1	-2
S3	-1	+1	0	-2
S4	-2	0	+1	+1

Initialize the Q' table with random values

- Choose an action a to perform in the current state, s
- Perform a and receive reward r
- Observe the new state, s'
- Update:

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha \left(\underbrace{R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a)}_{\text{learned value}} - Q_t(s_t, a_t) \right)$$

↑ **new state** ↑ **old state** ↑ **reward**

$s = s'$


- If the next state is not terminal, go back to step 1

Off-policy RL

Deep Q-Network, improving Q-Learning

Goal: learn an approximator $Q_\theta(s, a)$ for the optimal action-value function $Q^*(s, a)$

- Q-Learning suffer from
 - Memorization explosion
 - Instability and divergence when combined with a nonlinear Q-value and bootstrapping
- Deep Q-Network improvements
 - Convolutional layers: to consider compact / structural representations
 - Experience replay: storing episode steps into one replay memory D_t , and then randomly drawing batches of them to improve data efficiency and distribution
 - Separate target network: second network to generate the target Q-values that are periodically updated to compute the loss for every action during training



Initialize replay memory D to capacity N , action-value Q function with random weights θ , and target action-value \hat{Q} with weights $\theta^- = \theta$

1. Gather and store samples in a replay buffer with current policy
2. Random sample batches of experiences from the replay buffer
3. Use the sampled experiences to update the Q-network

$$L(\theta_i) = \mathbb{E}_{(s,a,r,s')}[(y_i - \overbrace{Q(s,a;\theta_i)}^{\text{prediction}})^2], \text{ where}$$
$$y_i = \begin{cases} R_T & \text{for terminal state } s_t \\ \underbrace{R_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a'; \theta^-)}_{\text{target}} & \text{for nonterminal } s_t \end{cases}$$

↑ reward ↓ decay

4. Repeat 1-3

General Conclusions

Model-based

- + 'Easy' to learn a model (supervised learning)
- + Learns *'all there is to know'* from the data
- Objective captures irrelevant information
- May focus on computing irrelevant details
- Computing policy (i.e. planning) is non-trivial
- Computing policy can be computationally expensive

Value-based

- + Closer to true objective
- + Fairly well understood
- Still not the true objective

Policy-based

- + Right objective
- + Better convergence properties
- + Effective in high-dimensional or continuous action spaces
- + Can learn stochastic policies
- Ignores other learnable knowledge
- Typically converge to a local rather than global optimum
- Evaluating a policy is normally inefficient and high variance is present

Challenges in Deep RL

Final thoughts and research directions

Core algorithms

- Stability – does your algorithm converge? (how many runs consistently)
 - Minimization of fitting error
 - Optimization on true objective
- Efficiency – how long does it take to converge? (how many samples)
 - Impact of off-policy vs. on-policy
- Generalization – after it converges, does it generalize?

Core assumptions

- Is this the right problem formulation? (better fit or more practical)
- What is the source of supervision? (efficient source)

Thank you!
Any question?

Eduardo Marques Pereira
edupereira@deloitte.pt